

```
In [112... #Load libraries
import os
#I Have problems with packages try running without this command
os.environ['KMP_DUPLICATE_LIB_OK']='True'
import numpy as np
import torch
import glob
import torch.nn as nn
import torch.nn.functional as F
from torchvision.transforms import transforms
from torch.utils.data import DataLoader
from torch.optim import Adam
from torch.autograd import Variable
import torchvision
import pathlib
import sys
import platform
# from google.colab import drive #used to access files
# drive.mount('/content/gdrive')
```

```
In [113... #checking for device
has_gpu = torch.cuda.is_available()
has_mps = torch.backends.mps.is_built()
device = "mps" if torch.backends.mps.is_built() \
    else "gpu" if torch.cuda.is_available() else "cpu"

print(f"Python Platform: {platform.platform()}")
print(f"PyTorch Version: {torch.__version__}")
print()
print(f"Python {sys.version}")
print(f"Pandas {pd.__version__}")
print(f"Scikit-Learn {sk.__version__}")
print("GPU is", "available" if has_gpu else "NOT AVAILABLE")
print("MPS (Apple Metal) is", "AVAILABLE" if has_mps else "NOT AVAILABLE")
print(f"Target device is {device}")
```

Python Platform: macOS-13.2-arm64-arm-64bit

PyTorch Version: 2.1.0.dev20230617

Python 3.9.16 | packaged by conda-forge | (main, Feb 1 2023, 21:38:11)

[Clang 14.0.6]

Pandas 2.0.2

Scikit-Learn 1.2.2

GPU is NOT AVAILABLE

MPS (Apple Metal) is AVAILABLE

Target device is mps

```
In [114... print(device)
```

mps

```
In [115... #Transforms
transformer=transforms.Compose([
    transforms.Resize((150,150)), #recommended size with 3x3 filters bigger
    transforms.RandomHorizontalFlip(), #chance of image being flipped is 0.5
    transforms.ToTensor(), #0-255 to 0-1, numpy to tensors (Pytorch uses te
    transforms.Normalize([0.5,0.5,0.5],
                        [0.5,0.5,0.5]) # 0-1 to [-1,1] , formula (x-mean)/s
])
```

```
In [116... #Dataloader

#Path for training and testing directory
train_path='/Users/hadi/Desktop/Concordia/Comp 6721/AIproject/fruits/trainin
test_path='/Users/hadi/Desktop/Concordia/Comp 6721/AIproject/fruits/testing/

train_loader=DataLoader(
    torchvision.datasets.ImageFolder(train_path,transform=transformer),
    batch_size=64, shuffle=True
)
#by default it was 64 but ta said use 32 (Personally I think 64 is better bec

test_loader=DataLoader(
    torchvision.datasets.ImageFolder(test_path,transform=transformer),
    batch_size=32, shuffle=True
)
```

```
In [117... #Returning classes
root=pathlib.Path(train_path)
classes = []
for j in root.iterdir():
    if j.name.split('/')[-1] != ".DS_Store":
        classes.append(j.name.split('/')[-1])
```

```
In [118... print(classes)

['Banana_Training', 'Kiwi_Training', 'Mango_Training', 'Orange_Training', 'Pl
um_Training', 'Apple_Training']
```

```
In [119... #Building CNN netwrok

class ConvNet(nn.Module):
    def __init__(self,num_classes=6):
        super(ConvNet,self).__init__()

        #Output size after convolution filter
        #((w-f+2P)/s) +1 w: width (32) - f: kernel size (filter) - p:padding
```

```
#Input shape= (32,3,32,32) 32: batch size - 3: rgb channels - 32x32:

self.conv1=nn.Conv2d(in_channels=3,out_channels=12,kernel_size=3,stroke=1)
#Shape= (256,12,150,150)
self.bn1=nn.BatchNorm2d(num_features=12)
#Shape= (256,12,150,150)
self.relu1=nn.ReLU()
#Shape= (256,12,150,150)

self.pool=nn.MaxPool2d(kernel_size=2)
#Reduce the image size by factor 2
#Shape= (256,12,75,75)

self.conv2=nn.Conv2d(in_channels=12,out_channels=20,kernel_size=3,stroke=1)
#Shape= (256,20,75,75)
self.relu2=nn.ReLU()
#Shape= (256,20,75,75)

self.conv3=nn.Conv2d(in_channels=20,out_channels=32,kernel_size=3,stroke=1)
#Shape= (256,32,75,75)
self.bn3=nn.BatchNorm2d(num_features=32)
#Shape= (256,32,75,75)
self.relu3=nn.ReLU()
#Shape= (256,32,75,75)

self.fc=nn.Linear(in_features=75 * 75 * 32,out_features=num_classes)

#Feed forward function

def forward(self,input):
    output=self.conv1(input)
    output=self.bn1(output)
    output=self.relu1(output)

    output=self.pool(output)

    output=self.conv2(output)
    output=self.relu2(output)

    output=self.conv3(output)
    output=self.bn3(output)
    output=self.relu3(output)
```

```

        #Above output will be in matrix form, with shape (256,32,75,75)

        output=output.view(-1,32*75*75)

        output=self.fc(output)

        return output

```

```
In [120...] model=ConvNet(num_classes=6).to(device)
```

```
In [121...] #Optimizer and loss function
optimizer=Adam(model.parameters(),lr=0.001,weight_decay=0.0001)
loss_function=nn.CrossEntropyLoss()
```

```
In [122...] num_epochs=10
```

```
In [123...] #calculating the size of training and testing images
train_count=len(glob.glob(train_path+'/**/*.png'))
test_count=len(glob.glob(test_path+'/**/*.png'))
print(train_count,test_count)
```

```
18341 4583
```

```
In [124...] #Model training and saving best model

best_accuracy=0.0
metrics = np.empty((0, 3)) # Initialize an empty matrix

for epoch in range(num_epochs):

    #Evaluation and training on training dataset
    model.train()
    train_accuracy=0.0
    train_loss=0.0

    for i, (images,labels) in enumerate(train_loader):
        if torch.backends.mps.is_built():
            images=Variable(images.to(device))
            labels=Variable(labels.to(device))

        optimizer.zero_grad()

        outputs=model(images)
        loss=loss_function(outputs,labels)
        loss.backward()
        optimizer.step()
```

```
train_loss+= loss.cpu().data*images.size(0)
_,prediction=torch.max(outputs.data,1)

train_accuracy+=int(torch.sum(prediction==labels.data))

train_accuracy=train_accuracy/train_count
train_loss=train_loss/train_count

# Evaluation on testing dataset
model.eval()

test_accuracy=0.0
for i, (images,labels) in enumerate(test_loader):
    if torch.backends.mps.is_built():
        images=Variable(images.to(device))
        labels=Variable(labels.to(device))

    outputs=model(images)
    _,prediction=torch.max(outputs.data,1)
    test_accuracy+=int(torch.sum(prediction==labels.data))

test_accuracy=test_accuracy/test_count

metric_per_epoch = np.array([train_loss,train_accuracy,test_accuracy])
metrics = np.vstack((metrics, metric_per_epoch))

print('Epoch: '+str(epoch)+' Train Loss: '+str(train_loss)+' Train Accur

#Save the best model
if test_accuracy>best_accuracy:
    torch.save(model.state_dict(),'best_checkpoint.model')
    best_accuracy=test_accuracy
```

```

Epoch: 0 Train Loss: tensor(2.6601) Train Accuracy: 0.8807589553459463 Test A
ccuracy: 0.8937377263801004
Epoch: 1 Train Loss: tensor(0.3625) Train Accuracy: 0.9627610272068045 Test A
ccuracy: 0.9524329042112154
Epoch: 2 Train Loss: tensor(0.1979) Train Accuracy: 0.9773185758682733 Test A
ccuracy: 0.9650883700632773
Epoch: 3 Train Loss: tensor(0.1190) Train Accuracy: 0.9841884302927867 Test A
ccuracy: 0.9144665066550295
Epoch: 4 Train Loss: tensor(0.1029) Train Accuracy: 0.9864238591134616 Test A
ccuracy: 0.9605062186340825
Epoch: 5 Train Loss: tensor(0.2189) Train Accuracy: 0.9811896843138324 Test A
ccuracy: 0.9570150556404102
Epoch: 6 Train Loss: tensor(0.0875) Train Accuracy: 0.9898042636715555 Test A
ccuracy: 0.9415230198559895
Epoch: 7 Train Loss: tensor(0.1495) Train Accuracy: 0.9859876778801592 Test A
ccuracy: 0.9696705214924721
Epoch: 8 Train Loss: tensor(0.1279) Train Accuracy: 0.9878414481216946 Test A
ccuracy: 0.9570150556404102
Epoch: 9 Train Loss: tensor(0.0488) Train Accuracy: 0.9936208494629518 Test A
ccuracy: 0.9692341261182631

```

```

In [125... import matplotlib.pyplot as plt


train_count_orange=len(glob.glob(train_path+'/Orange_Training/*.png'))
train_count_kiwi=len(glob.glob(train_path+'/Kiwi_Training/*.png'))
train_count_apple=len(glob.glob(train_path+'/Apple_Training/*.png'))
train_count_plum=len(glob.glob(train_path+'/Plum_Training/*.png'))
train_count_mango=len(glob.glob(train_path+'/Mango_Training/*.png'))
train_count_banana=len(glob.glob(train_path+'/Banana_Training/*.png'))

# Count the frequency of each label
label_counts = {"Orange":train_count_orange,"Kiwi":train_count_kiwi,"Apple":

# Extract the labels and frequencies
x = list(label_counts.keys())
y = list(label_counts.values())

# Plotting the graph
plt.bar(x, y)
plt.xlabel('Classes')
plt.ylabel('Frequency')
plt.title('Class Frequency')
plt.show()

```




```
In [126... #Plot for the loss curve

loss_values = []

for m in metrics:
    loss_values.append(m[0])

# Generate x-axis values for epochs
epochs = range(1, len(loss_values) + 1)

# Plotting the loss curve
plt.plot(epochs, loss_values, 'b-o')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss Curve')
plt.grid(True)
plt.show()
```



In [127... *#plot for the test vs training accuracy to check overfitting*

```
training_accuracy = []
test_accuracy = []

for m in metrics:
    training_accuracy.append(m[1])
    test_accuracy.append(m[2])

# Generate x-axis values for epochs
epochs = range(1, len(training_accuracy) + 1)

# Plotting the test vs training accuracy
plt.plot(epochs, training_accuracy, 'b-o', label='Training Accuracy')
plt.plot(epochs, test_accuracy, 'r-o', label='Test Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Test vs Training Accuracy')
plt.legend()
plt.grid(True)
plt.show()
```




In []: